

ST: Parallel Analysis Assignment #4

Kuan-Chih Lee

Instructor: Prof. Stratis Ioannidis

Question 1:

1a. to 1b.

$$1(a). \quad RSE(U, V) = \sum_{i,j} (u_i^T V_j - r_{ij})^2 + \lambda \sum_i \|u_i\|_2^2 + \mu \sum_j \|V_j\|_2^2$$

$$\therefore \nabla_{u_i} RSE(U, V) = 2 \sum_j (u_i^T V_j - r_{ij}) \cdot V_j + 2\lambda u_i$$

$$\nabla_{V_j} RSE(U, V) = 2 \sum_i (u_i^T V_j - r_{ij}) \cdot u_i + 2\mu V_j$$

1(b). If $d=1$, Prove $RSE(U, V)$ is not convex.

Solu.

If $d=1$, u_i and V_j is a scalar.

Therefore,

$$\begin{aligned} \phi(U, V) &= \min_{U, V} \sum_{i,j} (u_i \cdot V_j - r_{ij})^2 \\ &= \min \sum (r_{ij}^2 - 2u_i V_j r_{ij} + u_i^2 V_j^2) \end{aligned}$$

SGD case.

$$\Rightarrow \nabla \phi(U, V) = \begin{bmatrix} 2u_i V_j^2 - 2V_j r_{ij} \\ 2u_i^2 V_j - 2u_i r_{ij} \end{bmatrix} = \begin{bmatrix} \frac{\partial \phi}{\partial u_i} \\ \frac{\partial \phi}{\partial V_j} \end{bmatrix}$$

$$\Rightarrow \nabla^2 \phi(U, V) = \begin{bmatrix} 2V_j^2 & 4u_i V_j - 2r_{ij} \\ 4u_i V_j - 2r_{ij} & 2u_i^2 \end{bmatrix} = \begin{bmatrix} \frac{\partial^2 \phi}{\partial u_i^2} & \frac{\partial^2 \phi}{\partial u_i \partial V_j} \\ \frac{\partial^2 \phi}{\partial V_j \partial u_i} & \frac{\partial^2 \phi}{\partial V_j^2} \end{bmatrix}$$

$$\begin{aligned} \det(\nabla^2 \phi) &= 4u_i^2 V_j^2 - (16u_i^2 V_j^2 - 16u_i V_j r_{ij} + 4r_{ij}^2) \\ &= - (12u_i^2 V_j^2 - 16u_i V_j r_{ij} + 4r_{ij}^2) \end{aligned}$$

$$\therefore \forall u_i, V_j, r_{ij} \geq 0, \quad \nabla^2 \phi \leq 0 \Rightarrow \text{not convex} \quad \#$$

1c. to 2a.

1(c) No, the best example is saddle point.
 At saddle point, one of dimensions is local minimal,
 but another dimension will be local maximum.

1(d) Recall 1(b).

$$\nabla \phi(u, v) = \begin{bmatrix} 2u_1 v_j^* - 2v_j r_{ij} \\ 2u_i^2 v_j - 2u_i r_{ij} \end{bmatrix} = \begin{bmatrix} \frac{\partial \phi}{\partial u_i} \\ \frac{\partial \phi}{\partial v_j} \end{bmatrix}$$
 Case $\nabla \phi(u=0, v=0) = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$ # QED.

2(a) $RSE(u^k, v^k) = \sum_j S_{ij}^k^2$
 $\nabla_u RSE(u^k, v^k) = 2 \sum_j (S_{ij}^k v_j^k) + 2\lambda u_i$
 $\nabla_v RSE(u^k, v^k) = 2 \sum_i (S_{ij}^k u_i) + 2\mu v_j$

Question 2:

2b.

```
def predict(u, v):
    """ Given a user profile uprof and an item profile vpro

    Inputs are:
    -u: user profile, in the form of a numpy array
    -v: item profile, in the form of a numpy array

    The return value is
    - the inner product <u,v>
    """
    return u.T.dot(v)

def pred_diff(r, u, v):
    """ Given a rating, a user profile u and an item profile v

    Inputs are:
    -r: the rating a user gave to an item
    -u: user profile, in the form of a numpy array
    -v: item profile, in the form of a numpy array

    The return value is the difference
    - 差 = <u,v> - r
    """
    return predict(u, v) - r
```

```

def gradient_u(delta,u,v):
    """ Given a user profile u and an item profile v, and the
        difference  $\delta = \langle u, v \rangle - r$ 
        of the square error loss:
        
$$l(u,v) = (\langle u, v \rangle - r)^2$$

        Inputs are:
        -  $\delta$ : the difference  $\langle u, v \rangle - r$ 
        - u: user profile, in the form of a numpy array
        - v: item profile, in the form of a numpy array

        The return value is
        - The gradient w.r.t. u
    """
    return 2*delta*v

def gradient_v(delta,u,v):
    """ Given a user profile u and an item profile v, and the
        difference  $\delta = \langle u, v \rangle - r$ 
        of the square error loss:
        
$$l(u,v) = (\langle u, v \rangle - r)^2$$

        Inputs are:
        -  $\delta$ : the difference  $\langle u, v \rangle - r$ 
        - u: user profile, in the form of a numpy array
        - v: item profile, in the form of a numpy array

        The return value is
        - the gradient w.r.t. v
    """
    return 2*delta*u

```

2c.

```

def generateItemProfiles(R,d,seed,sparkContext,N):
    """ Generate the item profiles from rdd R and store them in an RDD containing tuples of the form
        (j,vj)
        where v is a random np.array of dimension d.

        The random uis are generated using normalVectorRDD(), a function in RandomRDDs.

        Inputs are:
        - R: an RDD that contains the ratings in (user, item, rating) form
        - d: the dimension of the user profiles
        - seed: a seed to be used for in generating the random vectors
        - sparkContext: a spark context
        - N: the number of partitions to be used during joins, etc.

        The return value is an RDD containing the item profiles
    """
    V = R.map(lambda (i,j,rij):j).distinct(numPartitions = N)
    numItems = V.count()
    randRDD = RandomRDDs.normalVectorRDD(sparkContext, numItems, d, numPartitions=N, seed=seed)
    V = V.zipWithIndex().map(swap)
    randRDD = randRDD.zipWithIndex().map(swap)
    return V.join(randRDD,numPartitions = N).values()

```


Why do we initialize the user and item profiles to random values? What would happen if we initialized all profiles to be zero vectors?

Ans:

If we initialize user and item profiles with zero vectors, the model will lose the capability of training. Recall question 2(b), the gradient, as well as regularization term, is the product of user or item profile.

Question 4:

4a.

```
def joinAndPredictAll(R,U,V,N):
    """ Receives as inputs the ratings R, the user profiles U, and the items V, and constructs a joined RDD.

    Inputs are:
    - R: an RDD containing tuples of the form (i,j,rij)
    - U: an RDD containing tuples of the form (i,ui)
    - V: an RDD containing tuples of the form (j,vj)
    - N: the number of partitions to be used during joins, etc.

    The output is a joined RDD containing tuples of the form:

    (i,j, 帮ij,ui,vj)

    where
    帮ij = <u,v>-rij
    is the prediction difference.

    """
    return R.map(lambda (i,j,rij): (i, (rij,j))).join(U, numPartitions = N) \
        .map(lambda (i, ((rij,j),ui)): (j, (i,ui,rij))).join(V, numPartitions = N) \
        .map(lambda (j, ((i,ui,rij),vj)): (i, j, pred_diff(rij,ui,vj), ui, vj))
```

4b.

```
def SE(joinedRDD):
    """ Receives as input a joined RDD as well as a 参 and a 参 and computes the MSE:

    SE(R,U,V) = 参_{(i,j in data)} (<ui,vj>-rij)^2

    The input is
    -joinedRDD: an RDD with tuples of the form (i,j, 帮ij,ui,vj), where 帮ij = <ui,vj> - rij is the prediction difference.

    The output is the SE.

    """
    return joinedRDD.map(lambda (i, j, delta, ui, vj): delta**2).reduce(add)

def normSqRDD(profileRDD,param):
    """ Receives as input an RDD of profiles (e.g., U) as well as a parameter (e.g., 参) and computes the square of norms:
    参 参_i ||ui||_2^2

    The input is:
    -profileRDD: an RDD of the form (i,u), where i is an index and u is a numpy array
    -param: a scalar 参>0

    The return value is:
    参 参_i ||ui||_2^2

    """
    return param * profileRDD.map(lambda (i, x): np.dot(x.T, x)).reduce(add)
```

SE is loss function, and normSqRDD is regularization term

4c.

```
def adaptU(joinedRDD,gamma,lam,N):
    """ Receives as input a joined RDD
    as well as a gain  $\gamma$ , and regularization parameters  $\lambda$  and  $\mu$ , and constructs a new RDD of user profiles of the form

    
$$u_i = u_i - \frac{\gamma}{N} \sum_j u_j \text{ RegSE}(R,U,V)$$


    where

    
$$\text{RegSE}(R,U,V) = \sum_{(i,j) \in R} (\langle u_i, v_j \rangle - r_{ij})^2 + \lambda \sum_i \|u_i\|_2^2 + \mu \sum_j \|v_j\|_2^2$$


    Inputs are
    -joinedRDD: an RDD with tuples of the form (i,j,rij,ui,vj), where  $rij = \langle u_i, v_j \rangle - r_{ij}$ 
    -gamma: the gain  $\gamma$ 
    -lam: the regularization parameter  $\lambda$ 
    -N: the number of partitions to be used in reduceByKey operations

    The return value is an RDD with tuples of the form (i,ui). The returned rdd contains exactly N partitions.
    """
    U_t0 = joinedRDD.map(lambda (i, j, delta, ui, vj): (i, ui)).reduceByKey(lambda x,y: x).repartition(N)
    grad = joinedRDD.map(lambda (i, j, delta, ui, vj): (i, gradient_u(delta,ui,vj))).reduceByKey(add).repartition(N)
    return U_t0.join(grad).map(lambda (i, (ui,g)): (i, ui - gamma*(g+2*lam*ui)))

def adaptV(joinedRDD,gamma,mu,N):
    """ Receives as input a joined RDD
    as well as a gain  $\gamma$ , and regularization parameters  $\lambda$  and  $\mu$ , and constructs a new RDD of user profiles of the form

    
$$u_i = u_i - \frac{\gamma}{N} \sum_j u_j \text{ RegSE}(R,U,V)$$


    where

    
$$\text{RegSE}(R,U,V) = \sum_{(i,j) \in R} (\langle u_i, v_j \rangle - r_{ij})^2 + \lambda \sum_i \|u_i\|_2^2 + \mu \sum_j \|v_j\|_2^2$$


    Inputs are
    -joinedRDD: an RDD with tuples of the form (i,j,rij,ui,vj), where  $rij = \langle u_i, v_j \rangle - r_{ij}$ 
    -gamma: the gain  $\gamma$ 
    -mu: the regularization parameter  $\mu$ 
    -N: the number of partitions to be used in reduceByKey operations

    The return value is an RDD with tuples of the form (j,vj). The returned rdd contains exactly N partitions.
    """
    V_t0 = joinedRDD.map(lambda (i, j, delta, ui, vj): (j, vj)).reduceByKey(lambda x,y: x).repartition(N)
    grad = joinedRDD.map(lambda (i, j, delta, ui, vj): (j, gradient_v(delta,ui,vj))).reduceByKey(add).repartition(N)
    return V_t0.join(grad).map(lambda (j, (vj,g)): (j, vj - gamma*(g+2*mu*vj)))
```

Question 5:

5a.

```
265 folds = {}
266
267 if args.output is None:
268     for k in range(args.folds):
269         folds[k] = readRatings(args.data+"/"+fold"+str(k),sc)
270 else:
271     folds[0] = readRatings(args.data,sc)
272
273 cross_val_rmse = []
274 for k in folds:
275     train_folds = [folds[j] for j in folds if j is not k]
276
277     if len(train_folds)>0:
278         train = train_folds[0]
279         for fold in train_folds[1:]:
280             train=train.union(fold)
281         train.repartition(args.N).cache()
282         test = folds[k].repartition(args.N).cache()
283         Mtrain=train.count()
284         Mtest=test.count()
285
286         print("Initiating fold %d with %d train samples and %d test samples" % (k,Mtrain,Mtest) )
287     else:
288         train = folds[k].repartition(args.N).cache()
289         test = train
290         Mtrain=train.count()
291         Mtest=test.count()
292         print("Running single training over training set with %d train samples. Test RMSE computes RMSE on training set" % Mtrain )
```

Data collection:

Line 267:

If args.output is not None, cross validation is skipped, and U,V are trained over entire dataset and store it in files output_U and output_V.'

Line 268-269

Read k in k-fold by args.folds. And then, read data from files under arg.data/ folder

and store in RDD in a dictionary called fold.

K-fold CV started (Line 274)

Line 275:

Store different folds, except the fold picked in Line 274, in a list called train_folds.

Line 279-280:

Pop out repeated data within different folds.

Line 281:

Training data will be the combination of k-1 folds

Testing data will be the fold picked at Line 274

5b.

Take question 5c as an example.

Original gamma is 0.001 (args.gain) and decay parameter is 0.2 (args.pow).

So, the first iteration adjusted gamma will be $\gamma / 1^{0.2}$.

Second iteration adjusted gamma will be $\gamma / 2^{0.2}$

20th iteration adjusted gamma will be $\gamma / 20^{0.2}$

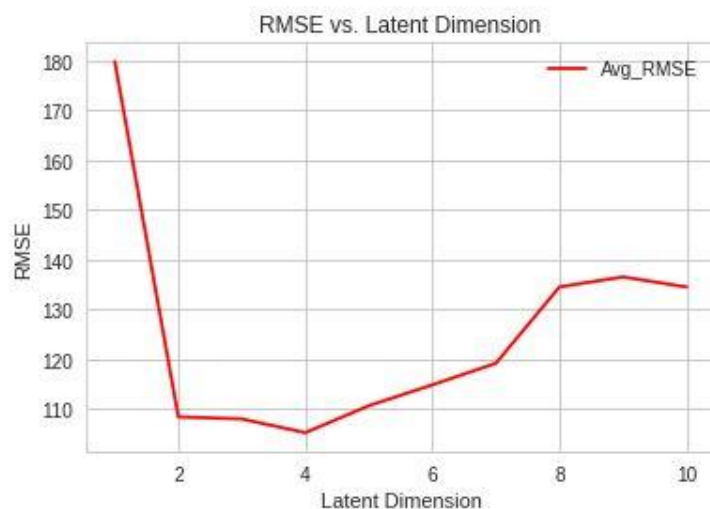
Formula:

$$\gamma \leftarrow \frac{\gamma}{\text{epoch}^{\text{pow}}}$$

5c.

Best Latent 4, Best Regul 0, RMSE 105.232654

```
df = pd.read_csv("A4_plotCSV_1", names=["Latent", "Regul", "Avg_RMSE"])
plt.plot(df.Latent, df.Avg_RMSE, c='r')
plt.xlabel('Latent Dimension')
plt.ylabel('RMSE')
plt.title('RMSE vs. Latent Dimension')
plt.legend()
plt.show()
```



5d.

It's not converged.

--gain 0.1 --pow 0.0 --maxiter 5

```
Initiating fold 4 with 40 train samples and 10 test samples
Training set contains 10 users and 9 items
Iteration: 1 Time: 1.557365 Objective: 1065416.059115 TestRMSE: 143.052712
Iteration: 2 Time: 4.907680 Objective: 4708796697.763937 TestRMSE: 4176.174975
Iteration: 3 Time: 8.185757 Objective: 71613325182300585916891136.000000 TestRMSE: 70913861720.276382
Iteration: 4 Time: 11.428958 Objective: 519500992543538657524979980811612278631390389010562588937419991766074916864.000000 TestRMSE: 683525349844142
64011839054033190912.000000
Iteration: 5 Time: 14.813233 Objective: 21887297067411361638785913257145970059031925886010820470233670906368437472242306836782415311051557864469568610
5408263574626883591136671306271601936705274963710699936582038202651267344966369852352058985732883566424410292224.000000 TestRMSE: 66188753211925384812877
4110560879852183068741501027648603949412404971274611759794754605785364630390147074816.000000
Latent 4, regularization 0, average error is: 12349339051303418524594015304431999426010013019140490722692760436573862985325091490057307126389656332613451
776.000000
```

Because of small and constant learning rate, it converges very slowly.

--gain 0.0001 --pow 1.0 --maxiter 5

```
Iteration: 1 Time: 1.797256 Objective: 1013406.331867 TestRMSE: 165.467171
Iteration: 2 Time: 5.374592 Objective: 1010664.098955 TestRMSE: 165.463864
Iteration: 3 Time: 9.070005 Objective: 1009317.129034 TestRMSE: 165.453203
Iteration: 4 Time: 12.596313 Objective: 1008410.193909 TestRMSE: 165.442764
Iteration: 5 Time: 15.780084 Objective: 1007720.662149 TestRMSE: 165.433156
Iteration: 6 Time: 19.033258 Objective: 1007161.333454 TestRMSE: 165.424353
Iteration: 7 Time: 22.474347 Objective: 1006688.979356 TestRMSE: 165.416245
Iteration: 8 Time: 25.869146 Objective: 1006278.981426 TestRMSE: 165.408726
Iteration: 9 Time: 29.121583 Objective: 1005915.964738 TestRMSE: 165.401710
Iteration: 10 Time: 32.425149 Objective: 1005589.671713 TestRMSE: 165.395127
Iteration: 11 Time: 35.812656 Objective: 1005292.910156 TestRMSE: 165.388918
Iteration: 12 Time: 39.320272 Objective: 1005020.437703 TestRMSE: 165.383038
Iteration: 13 Time: 42.554505 Objective: 1004768.312337 TestRMSE: 165.377449
Iteration: 14 Time: 46.000265 Objective: 1004533.493091 TestRMSE: 165.372118
Iteration: 15 Time: 49.216399 Objective: 1004313.583448 TestRMSE: 165.367019
Iteration: 16 Time: 52.451824 Objective: 1004106.660237 TestRMSE: 165.362130
Iteration: 17 Time: 55.668734 Objective: 1003911.155930 TestRMSE: 165.357431
Iteration: 18 Time: 59.002831 Objective: 1003725.775495 TestRMSE: 165.352905
Iteration: 19 Time: 62.320994 Objective: 1003549.436279 TestRMSE: 165.348538
Iteration: 20 Time: 65.534573 Objective: 1003381.223651 TestRMSE: 165.344318
```

5e.

The result shows that without any regularization, we can obtain the best average RMSE by cross-validation.

Best Latent 4, Best Regul 0, RMSE 108.866230

```
df = pd.read_csv("A4_plotCSV_2", names=["Latent", "Regul", "Avg_RMSE"])
plt.plot(df.Regul, df.Avg_RMSE, c='r')
plt.xlabel('Regularization')
plt.ylabel('RMSE')
plt.title('RMSE vs. Regularization')
plt.legend()
plt.show()
```

